

```
In [1]: # Import Necessary Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pandas_profiling #pip install pandas_profiling
import plotly.offline as po
import plotly.graph_objs as go
```

```
In [2]: #Perform Exploratory Data Analysis in just one line of code
pandas_profiling.ProfileReport(pd.read_csv('Tel_Customer_Churn_Dataset.csv'))

Summarize dataset:  0%|          | 0/5 [00:00<?, ?it/s]
Generate report structure:  0%|          | 0/1 [00:00<?, ?it/s]
Render HTML:  0%|          | 0/1 [00:00<?, ?it/s]
```

Overview

Dataset statistics

Number of variables	21
Number of observations	7043
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	1.1 MiB
Average record size in memory	168.0 B

Variable types

Categorical	14
Boolean	5
Numeric	2

Alerts

customerID has a high cardinality: 7043 distinct values	High cardinality
TotalCharges has a high cardinality: 6531 distinct values	High cardinality
tenure is highly overall correlated with Contract	High correlation

Out[2]:

```
In [3]: #Import Customer Churn Dataset
df = pd.read_csv('Tel_Customer_Churn_Dataset.csv')
df.head()
```

Out[3]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic

5 rows × 21 columns

In [4]: `df.shape`

Out[4]: (7043, 21)

In [5]:

```
# Convert String values (Yes and No) of Churn column to 1 and 0
df.loc[df.Churn=='No', 'Churn'] = 0
df.loc[df.Churn=='Yes', 'Churn'] = 1
```

In [6]:

```
# Convert 'No internet service' to 'No' for the below mentioned columns
cols = ['OnlineBackup', 'StreamingMovies', 'DeviceProtection',
        'TechSupport', 'OnlineSecurity', 'StreamingTV']
for i in cols :
    df[i] = df[i].replace({'No internet service' : 'No'})
```

In [7]:

```
# Replace all the spaces with null values
df['TotalCharges'] = df["TotalCharges"].replace(" ", np.nan)

# Drop null values of 'Total Charges' feature
df = df[df["TotalCharges"].notnull()]
df = df.reset_index()[df.columns]

# Convert 'Total Charges' column values to float data type
df["TotalCharges"] = df["TotalCharges"].astype(float)
df.head()
```

Out[7]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic

5 rows × 21 columns


In [8]: `df["Churn"].value_counts().values`

Out[8]: array([5163, 1869], dtype=int64)

```
In [9]: # Visualize Total Customer Churn
df_labels = df["Churn"].value_counts().keys().tolist()
df_values = df["Churn"].value_counts().values.tolist()

plot_data= [
    go.Pie(labels = df_labels,
           values = df_values,
           marker = dict(colors = [ 'Teal' , 'Grey'],
                          line = dict(color = "white",
                                      width = 1.5)),
           rotation = 90,
           hoverinfo = "label+value+text",
           hole = .6)
]
plot_layout = go.Layout(dict(title = "Customer Churn",
                              plot_bgcolor = "rgb(243,243,243)",
                              paper_bgcolor = "rgb(243,243,243)",))

fig = go.Figure(data=plot_data, layout=plot_layout)
po.iplot(fig)
```



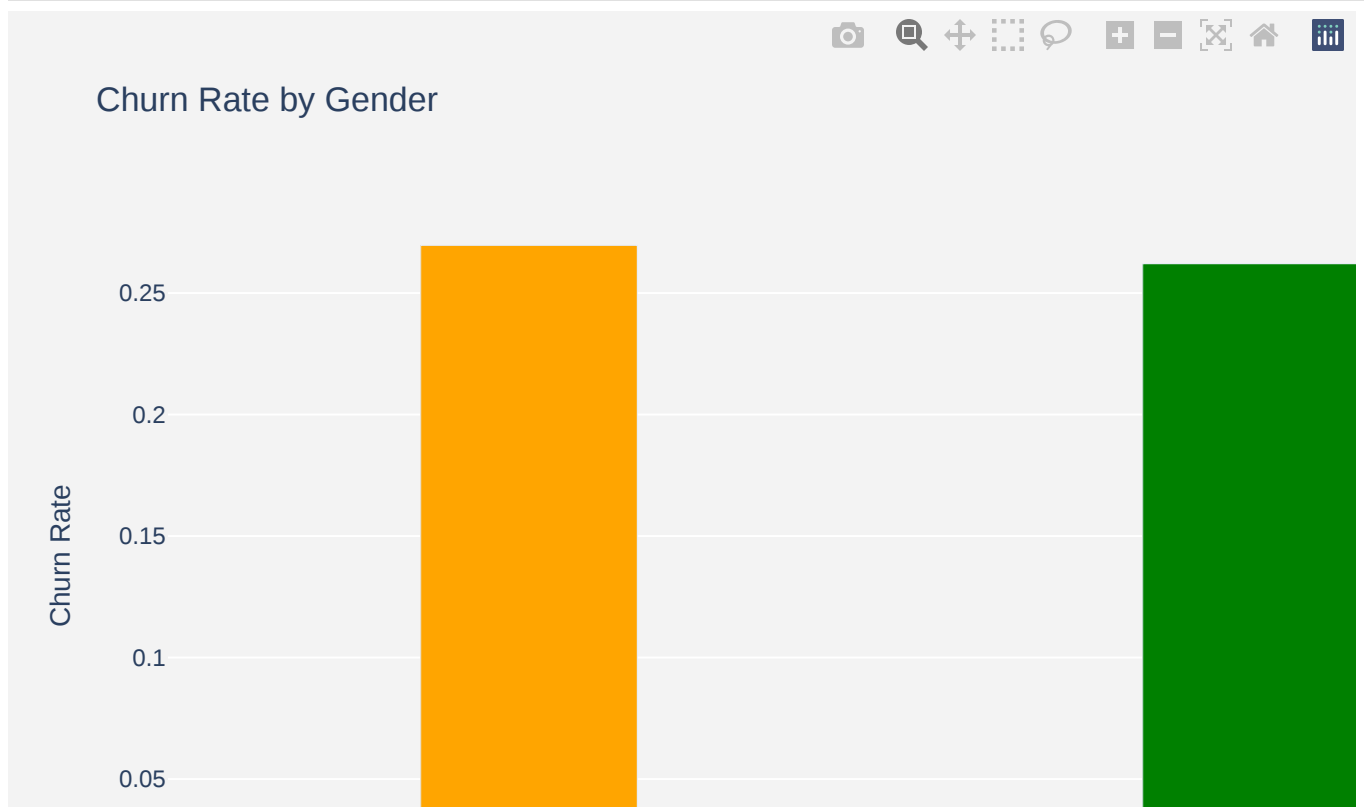
26.6% customer churn out while 73.4% customers stayed

```
In [10]: # Visualize Churn Rate by Gender
df_gender = df.groupby('gender').Churn.mean().reset_index()
```

```

go.Bar(
    x=df_gender['gender'],
    y=df_gender['Churn'],
    width = [0.3, 0.3],
    marker=dict(
        color=['orange', 'green'])
)
]
plot_layout = go.Layout(
    xaxis={"type": "category"},
    yaxis={"title": "Churn Rate"},
    title='Churn Rate by Gender',
    plot_bgcolor = 'rgb(243,243,243)',
    paper_bgcolor = 'rgb(243,243,243)',
)
fig = go.Figure(data=plot_data, layout=plot_layout)
po.iplot(fig)
print(df_gender)

```



```

gender    Churn
0  Female  0.269595
1   Male  0.262046

```

churn rate is slightly higher for females than males

```

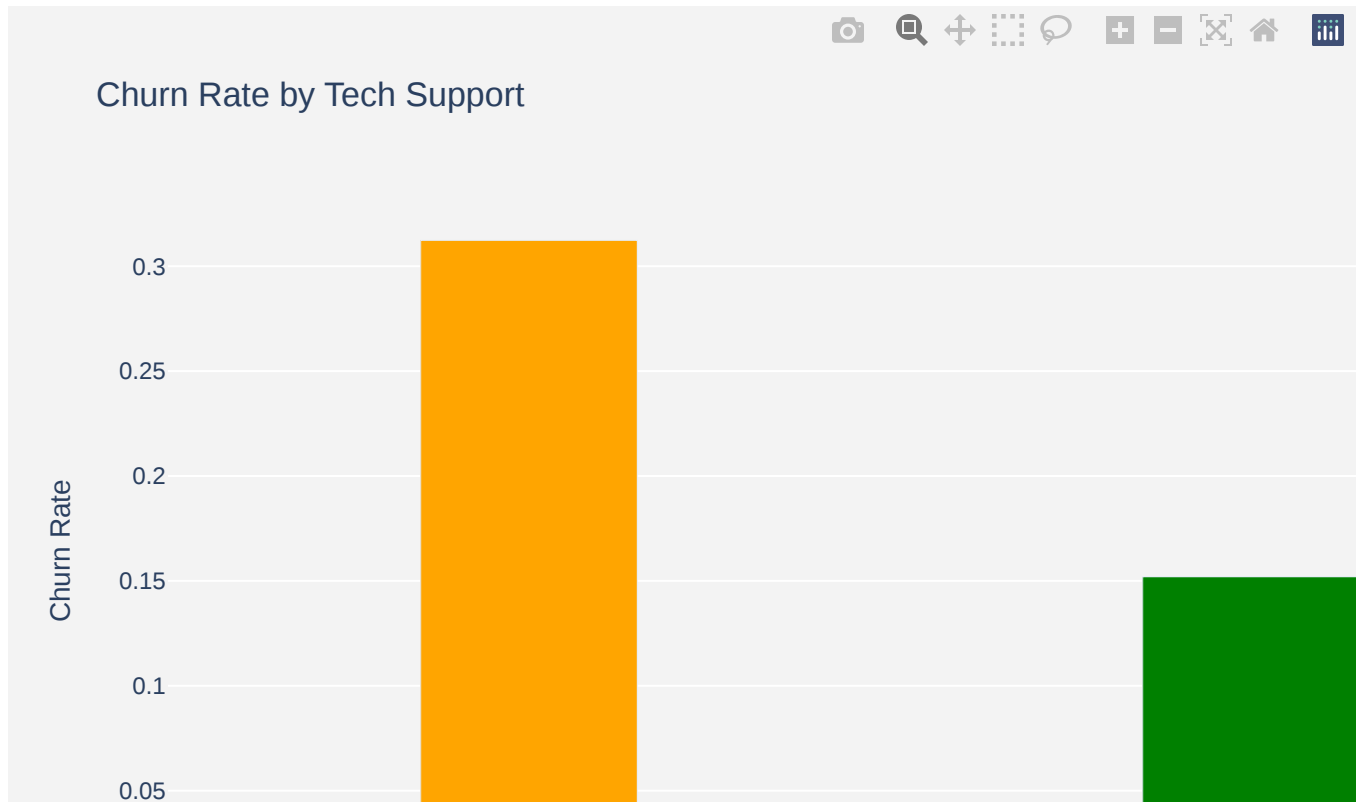
In [11]: # Visualize Churn Rate by Tech Support
df_techsupport = df.groupby('TechSupport').Churn.mean().reset_index()
plot_data = [
    go.Bar(
        x=df_techsupport['TechSupport'],
        y=df_techsupport['Churn'],
        width = [0.3, 0.3, 0.3],

```

```

        marker=dict(
            color=['orange', 'green', 'teal'])
    )
]
plot_layout = go.Layout(
    xaxis={"type": "category"},
    yaxis={"title": "Churn Rate"},
    title='Churn Rate by Tech Support',
    plot_bgcolor = 'rgb(243,243,243)',
    paper_bgcolor = 'rgb(243,243,243)',
)
fig = go.Figure(data=plot_data, layout=plot_layout)
po.iplot(fig)
print(df_techsupport)

```



	TechSupport	Churn
0	No	0.312300
1	Yes	0.151961

churn rate for those with no tech support is higher compared to those without tech support

```


In [12]: # Visualize Churn Rate by Internet Services
df_internet_service = df.groupby('InternetService').Churn.mean().reset_index()
plot_data = [
    go.Bar(
        x=df_internet_service['InternetService'],
        y=df_internet_service['Churn'],
        width = [0.3, 0.3, 0.3],
        marker=dict(
            color=['orange', 'green', 'teal'])
    )
]

```

```

plot_layout = go.Layout(
    xaxis={"type": "category"},
    yaxis={"title": "Churn Rate"},
    title='Churn Rate by Internet Service',
    plot_bgcolor = 'rgb(243,243,243)',
    paper_bgcolor = 'rgb(243,243,243)',
)
fig = go.Figure(data=plot_data, layout=plot_layout)
po.iplot(fig)
print(df_internet_service)

```



```

InternetService    Churn
0          DSL    0.189983
1    Fiber optic    0.418928
2             No    0.074342

```

we have alot more people churning on fiber optics compared to DSL

```


In [13]: # Visualize Churn Rate by Payment Method
df_payment = df.groupby('PaymentMethod').Churn.mean().reset_index()
plot_data = [
    go.Bar(
        x=df_payment['PaymentMethod'],
        y=df_payment['Churn'],
        width = [0.3, 0.3,0.3,0.3],
        marker=dict(
            color=['orange', 'green', 'teal', 'magenta'])
    )
]
plot_layout = go.Layout(
    xaxis={"type": "category"},
    yaxis={"title": "Churn Rate"},

```

```

        title='Churn Rate by Payment Method',
        plot_bgcolor = 'rgb(243,243,243)',
        paper_bgcolor = 'rgb(243,243,243)',
    )
fig = go.Figure(data=plot_data, layout=plot_layout)
po.iplot(fig)
print(df_payment)

```



```

      PaymentMethod    Churn
0  Bank transfer (automatic)  0.167315
1   Credit card (automatic)  0.152531
2      Electronic check    0.452854
3      Mailed check        0.192020

```

we find more people who make payment through electronic check churning compared to the other payment methods

```

In [14]: # Visualize Churn Rate by Contract Duration
df_contract = df.groupby('Contract').Churn.mean().reset_index()
plot_data = [
    go.Bar(
        x=df_contract['Contract'],
        y=df_contract['Churn'],
        width = [0.3, 0.3, 0.3],
        marker=dict(
            color=['orange', 'green', 'teal'])
    )
]
plot_layout = go.Layout(
    xaxis={"type": "category"},
    yaxis={"title": "Churn Rate"},


```



```

        title='Churn Rate by Contract Duration',
        plot_bgcolor = 'rgb(243,243,243)',
        paper_bgcolor = 'rgb(243,243,243)',
    )
fig = go.Figure(data=plot_data, layout=plot_layout)
po.iplot(fig)
print(df_contract)

```



```

      Contract      Churn
0  Month-to-month  0.427097
1      One year    0.112772
2      Two year    0.028487

```

we have more people churning on the month to month contract compared to the one and two years contracts

```


In [15]: # Visualize Relation between Tenure & Churn rate
df_tenure = df.groupby('tenure').Churn.mean().reset_index()
plot_data = [
    go.Scatter(
        x=df_tenure['tenure'],
        y=df_tenure['Churn'],
        mode='markers',
        name='Low',
        marker= dict(size= 5,
                    line= dict(width=0.8),
                    color= 'green'
                    ),
    ),
]
plot_layout = go.Layout(

```

```

    yaxis= {'title': "Churn Rate"},
    xaxis= {'title': "Tenure"},
    title='Relation between Tenure & Churn rate',
    plot_bgcolor = "rgb(243,243,243)",
    paper_bgcolor = "rgb(243,243,243)",
)
fig = go.Figure(data=plot_data, layout=plot_layout)
po.iplot(fig)
print(df_tenure)

```



```

    tenure    Churn
0          1  0.619902
1          2  0.516807
2          3  0.470000
3          4  0.471591
4          5  0.481203
..      ...      ...
67        68  0.090000
68        69  0.084211
69        70  0.092437
70        71  0.035294
71        72  0.016575

```

[72 rows x 2 columns]

there is a negative relationship between tenure and churn rate, as tenure increase churn rate falls

Data Preprocessing

In [16]: *#Perform One Hot Encoding using get_dummies method*

Loading [MathJax]/extensions/Safe.js dummies(df, columns = ['Contract', 'Dependents', 'DeviceProtection', 'gender',

```
'InternetService', 'MultipleLines',  
'OnlineSecurity', 'PaperlessBilli',  
'PaymentMethod', 'PhoneService', '  
'StreamingMovies', 'StreamingTV',
```

```
drop_first=True)
```

```
df.columns
```

```
Out[16]: Index(['customerID', 'tenure', 'MonthlyCharges', 'TotalCharges', 'Churn',  
            'Contract_One year', 'Contract_Two year', 'Dependents_Yes',  
            'DeviceProtection_Yes', 'gender_Male', 'InternetService_Fiber optic',  
            'InternetService_No', 'MultipleLines_No phone service',  
            'MultipleLines_Yes', 'OnlineBackup_Yes', 'OnlineSecurity_Yes',  
            'PaperlessBilling_Yes', 'Partner_Yes',  
            'PaymentMethod_Credit card (automatic)',  
            'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check',  
            'PhoneService_Yes', 'SeniorCitizen_1', 'StreamingMovies_Yes',  
            'StreamingTV_Yes', 'TechSupport_Yes'],  
            dtype='object')
```

```
In [17]: df.head()
```

```
Out[17]:
```

	customerID	tenure	MonthlyCharges	TotalCharges	Churn	Contract_One year	Contract_Two year	Dependents_Yes	D
0	7590-VHVEG	1	29.85	29.85	0	0	0	0	
1	5575-GNVDE	34	56.95	1889.50	0	1	0	0	
2	3668-QPYBK	2	53.85	108.15	1	0	0	0	
3	7795-CFOCW	45	42.30	1840.75	0	1	0	0	
4	9237-HQITU	2	70.70	151.65	1	0	0	0	

5 rows × 26 columns

```
In [18]: #Perform Feature Scaling and One Hot Encoding  
from sklearn.preprocessing import StandardScaler  
  
#Perform Feature Scaling on 'tenure', 'MonthlyCharges', 'TotalCharges' in order to bring  
standardScaler = StandardScaler()  
columns_for_ft_scaling = ['tenure', 'MonthlyCharges', 'TotalCharges']  
  
#Apply the feature scaling operation on dataset using fit_transform() method  
df[columns_for_ft_scaling] = standardScaler.fit_transform(df[columns_for_ft_scaling])  
df.head()
```

Out[18]:

	customerID	tenure	MonthlyCharges	TotalCharges	Churn	Contract_One year	Contract_Two year	Dependents_Yes
0	7590-VHVEG	-1.280248	-1.161694	-0.994194	0	0	0	0
1	5575-GNVDE	0.064303	-0.260878	-0.173740	0	1	0	0
2	3668-QPYBK	-1.239504	-0.363923	-0.959649	1	0	0	0
3	7795-CFOCW	0.512486	-0.747850	-0.195248	0	1	0	0
4	9237-HQITU	-1.239504	0.196178	-0.940457	1	0	0	0

5 rows × 26 columns

In [19]: `df.columns`

Out[19]: Index(['customerID', 'tenure', 'MonthlyCharges', 'TotalCharges', 'Churn', 'Contract_One year', 'Contract_Two year', 'Dependents_Yes', 'DeviceProtection_Yes', 'gender_Male', 'InternetService_Fiber optic', 'InternetService_No', 'MultipleLines_No phone service', 'MultipleLines_Yes', 'OnlineBackup_Yes', 'OnlineSecurity_Yes', 'PaperlessBilling_Yes', 'Partner_Yes', 'PaymentMethod_Credit card (automatic)', 'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check', 'PhoneService_Yes', 'SeniorCitizen_1', 'StreamingMovies_Yes', 'StreamingTV_Yes', 'TechSupport_Yes'], dtype='object')

In [20]: *#Create Feature variable X and Target variable y*

```
X = df.drop(['Churn','customerID'], axis = 1)
y = df['Churn'].astype("int")
```

In [21]: *#Split the data into training set (70%) and test set (30%)*

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_state
```

In [22]: *# Machine Learning classification model libraries*

```
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics
```

starting with logistic Regression model (used for binary classification)

In [23]: *#Fit the logistic Regression Model*

```
log = LogisticRegression(random_state=50).fit(X_train,y_train)
```

#Predict the value for new, unseen data

```
pred = log.predict(X_test)
```

Find Accuracy using accuracy_score method

```
logmodel_accuracy = round(metrics.accuracy_score(y_test, pred) * 100, 2)
print("R-square of Logistic regression: ",logmodel_accuracy)
```

support vector machine model (svm)

```
In [24]: #Fit the Support Vector Machine Model
svcmodel = SVC(kernel='linear', random_state=50, probability=True)
svcmodel.fit(X_train,y_train)

#Predict the value for new, unseen data
svc_pred = svcmodel.predict(X_test)

# Find Accuracy using accuracy_score method
svc_accuracy = round(metrics.accuracy_score(y_test, svc_pred) * 100, 2)
print("R-square of SVC: ",svc_accuracy)
```

R-square of SVC: 80.66

K-nearest Neighbour model

```
In [25]: #Fit the K-Nearest Neighbor Model
from sklearn.neighbors import KNeighborsClassifier
knnmodel = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2) #p=2 represents
knnmodel.fit(X_train, y_train)

#Predict the value for new, unseen data
knn_pred = knnmodel.predict(X_test)

# Find Accuracy using accuracy_score method
knn_accuracy = round(metrics.accuracy_score(y_test, knn_pred) * 100, 2)
print("R-Squared for KNN: ",knn_accuracy)
```

2023-02-21 16:12:25,780 [14492] WARNING py.warnings:109: [JupyterRequire] C:\Users\ADOW UONA-OW00\Anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning:

Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

R-Squared for KNN: 76.82

Decision tree classification model

```
In [26]: #Fit the Decision Tree Classification Model
from sklearn.tree import DecisionTreeClassifier
dtmodel = DecisionTreeClassifier(criterion = "gini", random_state = 50)
dtmodel.fit(X_train, y_train)

#Predict the value for new, unseen data
dt_pred = dtmodel.predict(X_test)

# Find Accuracy using accuracy_score method
dt_accuracy = round(metrics.accuracy_score(y_test, dt_pred) * 100, 2)
print("R-Squared for Decision tree: ", dt_accuracy)
```

R-Squared for Decision tree: 73.27

Random Forest Classification model

```
In [27]: #Fit the Random Forest Classification Model
from sklearn.ensemble import RandomForestClassifier
rfmodel = RandomForestClassifier(n_estimators = 100, criterion = 'entropy', random_state
rfmodel.fit(X_train, y_train)

#Predict the value for new, unseen data
rf_pred = rfmodel.predict(X_test)

# Find Accuracy using accuracy_score method
rf_accuracy = round(metrics.accuracy_score(y_test, rf_pred) * 100, 2)
print("R-square for Random forest: ", rf_accuracy)
```

R-square for Random forest: 79.38

```
In [28]: # Compare Several models according to their Accuracies
Model_Comparison = pd.DataFrame({
    'Model': ['Logistic Regression', 'Support Vector Machine', 'K-Nearest Neighbor',
             'Decision Tree', 'Random Forest'],
    'Score': [logmodel_accuracy, svc_accuracy, knn_accuracy,
             dt_accuracy, rf_accuracy]})
Model_Comparison_df = Model_Comparison.sort_values(by='Score', ascending=False)
Model_Comparison_df = Model_Comparison_df.set_index('Score')
Model_Comparison_df.reset_index()
```

```
Out[28]:
```

	Score	Model
0	81.14	Logistic Regression
1	80.66	Support Vector Machine
2	79.38	Random Forest
3	76.82	K-Nearest Neighbor
4	73.27	Decision Tree

Logistic regression has the highest accuracy score hence we compute the confusion matrix

```
In [29]: #Generate confusion matrix for logistics regression model as it has maximum Accuracy
from sklearn.metrics import confusion_matrix
conf_mat_logmodel = confusion_matrix(y_test, pred)
conf_mat_logmodel
```

```
Out[29]: array([[1396, 165],
                [ 233, 316]], dtype=int64)
```

1396 and 316 are the correct predictions and 233 and 165 are incorrect predictions

```
In [31]: # Predict the probability of Churn of each customer
df['Probability_of_Churn'] = log.predict_proba(df[X_test.columns])[:,1]
df.head(10)
```

Out[31]:

	customerID	tenure	MonthlyCharges	TotalCharges	Churn	Contract_One year	Contract_Two year	Dependents_Yes
0	7590-VHVEG	-1.280248	-1.161694	-0.994194	0	0	0	0
1	5575-GNVDE	0.064303	-0.260878	-0.173740	0	1	0	0
2	3668-QPYBK	-1.239504	-0.363923	-0.959649	1	0	0	0
3	7795-CFOCW	0.512486	-0.747850	-0.195248	0	1	0	0
4	9237-HQITU	-1.239504	0.196178	-0.940457	1	0	0	0
5	9305-CDSKC	-0.995040	1.158489	-0.645369	1	0	0	0
6	1452-KIOVK	-0.424625	0.807802	-0.147313	0	0	0	1
7	6713-OKOMC	-0.913552	-1.165018	-0.874169	0	0	0	0
8	7892-POOKP	-0.180161	1.329677	0.336516	1	0	0	0
9	6388-TABGU	1.205134	-0.287470	0.531476	0	1	0	1

10 rows × 27 columns

In [32]:

```
# Create a Dataframe showcasing probability of Churn of each customer
df[['customerID', 'Probability_of_Churn']].head(10)
```

Out[32]:

	customerID	Probability_of_Churn
0	7590-VHVEG	0.649225
1	5575-GNVDE	0.043673
2	3668-QPYBK	0.340977
3	7795-CFOCW	0.026396
4	9237-HQITU	0.694569
5	9305-CDSKC	0.782005
6	1452-KIOVK	0.490814
7	6713-OKOMC	0.290564
8	7892-POOKP	0.594619
9	6388-TABGU	0.011939

for probability of churn that is close to 1 means the customer may leave soon so we need to take corrective actions for such customers.

In []: